

Retrieval Experiment: Dense vs Hybrid+Rerank

Overview

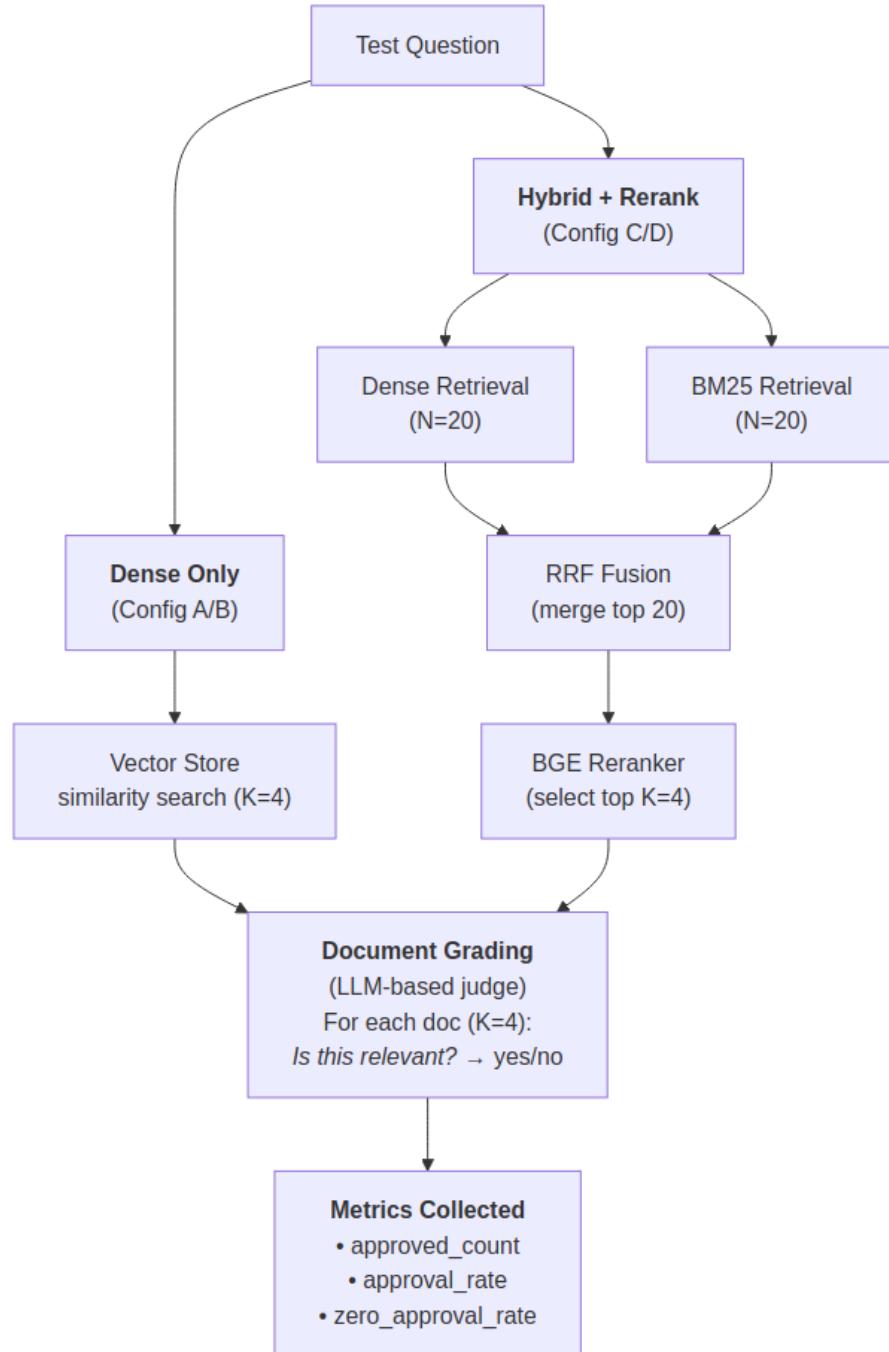
This experiment compares dense-only retrieval against a hybrid pipeline (Dense + BM25 + RRF + Rerank) for Case-Based Reasoning on mathematical problem solving. Retrieval quality is evaluated using an LLM-based document grader that judges whether retrieved examples contain transferable problem-solving techniques.

Experimental Setup

Configurations

Config	Dense Embeddings	BM25	RRF Fusion	Reranker
A. Dense (Nomic)	Nomic-embed-v1.5	-	-	-
B. Dense (Qwen3)	Qwen3-Embedding-8B	-	-	-
C. Hybrid+Rerank (Nomic)	Nomic-embed-v1.5	Yes	Yes	BGE-reranker-v2-m3
D. Hybrid+Rerank (Qwen3)	Qwen3-Embedding-8B	Yes	Yes	BGE-reranker-v2-m3

Pipeline



The hybrid pipeline performs parallel dense and BM25 retrieval (N=20 each), fuses results via Reciprocal Rank Fusion (RRF, k=60), then reranks with a cross-encoder to select the top K=4 documents.

Datasets

Dataset	Test Queries	Retrieval Corpus (Training Set)
GSM8K	1,319	7,473
MATH	5,000	7,500

Technical Components

Component	Implementation
Dense Embeddings	Nomic-embed-v1.5 (137M params) or Qwen3-Embedding-8B
Vector Store	ChromaDB
BM25	rank_bm25.BM25Okapi
Reranker	BAAI/bge-reranker-v2-m3 (CrossEncoder)
Grader LLM	Qwen2.5-32B-Instruct via vLLM

Documents are graded in a single batched vLLM call. The grader outputs a binary “yes”/“no” judgment on whether each document provides adaptable problem-solving techniques for the query.

Hyperparameters

Parameter	Value	Description
N_INITIAL	20	Documents retrieved per method
N_RERANK	20	Documents passed to reranker after RRF
TOP_K	4	Final documents graded
RRF_K	60	RRF smoothing constant

Results

GSM8K (1,319 queries)

Rank	Config	Avg Approved	Approval Rate	Zero Rate	Coverage
1	B. Dense (Qwen3)	3.40	85.0%	1.4%	98.6%
2	D. Hybrid+Rerank (Qwen3)	3.16	79.0%	3.0%	97.0%
3	A. Dense (Nomic)	3.14	78.5%	2.4%	97.6%
4	C. Hybrid+Rerank (Nomic)	3.11	77.7%	3.0%	97.0%

Approval Distribution (GSM8K)

Config	0 docs	1 doc	2 docs	3 docs	4 docs
Dense (Qwen3)	18	58	120	304	819
Hybrid+Rerank (Qwen3)	39	92	163	349	676
Dense (Nomic)	32	117	152	351	667
Hybrid+Rerank (Nomic)	40	105	173	355	646

MATH (5,000 queries)

Rank	Config	Avg Approved	Approval Rate	Zero Rate	Coverage
1	B. Dense (Qwen3)	3.48	86.9%	1.5%	98.5%
2	A. Dense (Nomic)	2.86	71.4%	6.1%	93.9%
3	D. Hybrid+Rerank (Qwen3)	2.79	69.8%	5.7%	94.3%
4	C. Hybrid+Rerank (Nomic)	2.63	65.8%	7.7%	92.3%

Coverage (queries with at least 1 relevant example)

Config	GSM8K	MATH
Dense (Qwen3)	98.6% (1301/1319)	98.5% (4927/5000)
Dense (Nomic)	97.6% (1287/1319)	93.9% (4697/5000)
Hybrid+Rerank (Qwen3)	97.0% (1280/1319)	94.3% (4716/5000)
Hybrid+Rerank (Nomic)	97.0% (1279/1319)	92.3% (4614/5000)

Analysis

BM25 Impact

Dataset	Dense Qwen3	Hybrid Qwen3	Penalty
GSM8K	85.0%	79.0%	-6.0pp
MATH	86.9%	69.8%	-17.1pp

Dataset	Dense Nomic	Hybrid Nomic	Penalty
GSM8K	78.5%	77.7%	-0.8pp
MATH	71.4%	65.8%	-5.6pp

The BM25 penalty increases with dataset complexity. On the more diverse MATH dataset, adding BM25 to Dense Qwen3 reduces approval rate by 17.1 percentage points.

Hyperparameter Sweep (GSM8K, Hybrid Qwen3)

N_INITIAL	N_RERANK	TOP_K	Approval	Avg Approved	Zero Rate
10	10	4	79.9%	3.20	2.4%
50	20	4	79.7%	3.19	2.7%
20 (baseline)	20	4	79.0%	3.16	3.0%
100	50	4	78.4%	3.13	2.7%
50	50	4	78.2%	3.13	2.5%

Smaller retrieval pools outperform larger ones. Expanding the candidate pool introduces more noise from BM25 that the reranker fails to filter.

Zero-Approval Query Analysis (Dense Nomic, GSM8K)

32 queries (2.4%) received zero approved documents. Breakdown by problem type: - 31% ratio/proportion problems - 22% rate/time problems - 22% multi-step problems

Root cause: embeddings match on topic similarity rather than mathematical reasoning structure. A “bakery” problem about unit conversion retrieves other “bakery” problems about quantity counting, not structurally similar unit conversion problems from other domains.

Key Findings

1. **Dense Qwen3 dominates** across both datasets (85-87% approval, 98.5%+ coverage).
2. **Dense outperforms Hybrid** for both embedding models on both datasets.
3. **Embedding model quality is the dominant factor.** Upgrading from Nomic (137M) to Qwen3 (8B) yields +6.5pp on GSM8K and +15.5pp on MATH — far exceeding any pipeline-level improvement.
4. **BM25 degrades retrieval quality for math.** Keyword matching surfaces topically similar but structurally different problems. The penalty grows with dataset complexity (-6pp GSM8K, -17pp MATH).
5. **The BGE reranker cannot compensate** for BM25 noise injection, as it lacks math reasoning awareness.
6. **These results contradict Anthropic’s Contextual Retrieval findings** (49-67% failure rate reduction from hybrid retrieval), suggesting that domain-specific retrieval tasks — particularly mathematical reasoning — do not benefit from generic hybrid pipelines.

Recommendation

Use **Dense Qwen3** for the CBR retrieval pipeline. It is the simplest configuration and outperforms all alternatives. The hybrid approach adds complexity while degrading quality.